

Assignment 11

Task 1

Snippet a

c1 is of type unsigned

```
unsigned c2 = 32 * c1;
```

If we assume that the multiplication takes longer than a bit shift, we can optimize this code as follows:

```
unsigned c2 = c1 << 5;
```

If we look at the compiled assembly code, we see no difference between those two snippets when compiled with -O3.

Snippet b

c1 is of type unsigned

```
unsigned c2 = 15 * c1;
```

Assuming bit shifting + one subtraction are faster than a multiplication, one could do this:

```
unsigned c2 = (c1 << 4) - c1;
```

Once again the corresponding assembly codes are identical.

Snippet c

c1 is of type unsigned

```
unsigned c2 = 96 * c1;
```

Assuming that 2 shift operations and one addition are still faster than a single multiplication, the following optimization could be helpful:

```
unsigned c2 = (c1 << 5) + (c1 << 6);
```

The compiler **disagrees** with this optimization. The 'optimized' snippet gets replaced by the following code:

```

c_solution:
    imul    esi, edi, 96
    xor     eax, eax
    mov     edi, OFFSET FLAT:.LC2
    jmp     printf

```

which basically performs a multiplication by 96, which means that the compiler thinks that the unoptimized version is faster than the optimized one.

Additionally, if we compile the unoptimized version, the compiler does the following:

```

c:
    lea     esi, [rdi+rdi*2]
    xor     eax, eax
    mov     edi, OFFSET FLAT:.LC2
    sal     esi, 5
    jmp     printf

```

Since we are not quite talented with x86 assembly code, we cannot explain this snippet.

Snippet d

`c1` is of type `unsigned`

```

unsigned c2 = 0.125 * c1;

```

Since the multiplication of 0.125 is equivalent to the division by 8, we can optimize this snippet as follows:

```

unsigned c2 = c1 >> 3;

```

The compiler translates this to the expected assembly code:

```

d_solution:
    mov     esi, edi
    xor     eax, eax
    mov     edi, OFFSET FLAT:.LC4
    shr     esi, 3
    jmp     printf

```

But if we look at the compiled unoptimized version, the compiler thinks this magic looking code is faster than a normal multiplication:

```

d:
    mov     edi, edi
    pxor    xmm0, xmm0
    xor     eax, eax
    cvtsi2sdq    xmm0, rdi
    mulsd    xmm0, QWORD PTR .LC3[rip]
    mov     edi, OFFSET FLAT:.LC4
    cvtt2sd2si    rsi, xmm0
    jmp     printf

```

Snippet e

`a` is of type `unsigned *`

```
unsigned sum_fifth = 0;
for (int i = 0; i < N / 5; ++i) {
    sum_fifth += a[5 * i];
}
```

If we want to get rid of the expensive multiplication inside the loop, we can simply transform the loop head like this:

```
unsigned sum_fifth = 0;
for (int i = 0; i < N; i+=5) {
    sum_fifth += a[i];
}
```

The compiled versions of the unoptimized and optimized snippet look identical.

Snippet f

`a` is of type `double *`

```
for (int i = 0; i < N; ++i) {
    a[i] += i / 5.3;
}
```

In this case, we want to get rid of the division with the float number `5.3` to optimize the code snippet. To do that, we could hard code the value of the constant `1/5.3` and convert it to a multiplication. But a multiplication in a for loop can be converted in to additions which is more efficient. Although `1/5.3` is an irrational number, if we implement this constant with enough decimal points, the accuracy might not be affected by these changes.

```
float h = 0.0;
for (int i = 0; i < N; ++i) {
    a[i] += h;
    h = h + 0.18867924528301886792452830188679;
}
```

The unoptimized version compiles to this:

```
movdqa xmm2, XMMWORD PTR .LC6[rip]
movdqa xmm4, XMMWORD PTR .LC7[rip]
lea rax, [rdi+8000]
movapd xmm3, XMMWORD PTR .LC8[rip]
.L17:
pshufd xmm0, xmm2, 238
cvtdq2pd xmm1, xmm2
movupd xmm6, XMMWORD PTR [rdi]
add rdi, 32
cvtdq2pd xmm0, xmm0
divpd xmm1, xmm3
movupd xmm5, XMMWORD PTR [rdi-16]
```

```

paddb    xmm2, xmm4
divpd    xmm0, xmm3
addpd    xmm1, xmm6
movups   XMMWORD PTR [rdi-32], xmm1
addpd    xmm0, xmm5
movups   XMMWORD PTR [rdi-16], xmm0
cmp      rax, rdi
jne      .L17
ret

```

While the optimized version this assembly code causes:

```

movsd    xmm2, QWORD PTR .LC10[rip]
lea      rax, [rdi+8000]
pxor     xmm0, xmm0
.L20:
movsd    xmm1, QWORD PTR [rdi]
cvtss2sd        xmm0, xmm0
add      rdi, 8
addsd    xmm1, xmm0
addsd    xmm0, xmm2
movsd    QWORD PTR [rdi-8], xmm1
cvtsd2ss        xmm0, xmm0
cmp      rax, rdi
jne      .L20
ret

```

Snippet g

c1 is of type float

```
float c2 = -1 * c1;
```

To just swap the sign, the multiplication of -1 is not the best way. Better is to just put - before the variable you want the negation of because then a `sign-bit flip` is executed. So most significant bit, which equals to the sign bit (IEEE 754 single-precision binary floating-point format). So in fact, only a XOR-operation is executed to negate a number.

```
float c2 = -c1;
```

The compiler does this by default so both result in the same assembly code:

```

xor      eax, eax
ret

```